

```
In [1]: import numpy as np
import pandas as pd
import random
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix,
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from sklearn.neural_network import MLPClassifier
import pyswarms as ps
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, GradientBoostingClassifier
from sklearn.metrics import roc_curve, auc
```

C:\Users\HP\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.8.4' or newer of 'numexpr' (version '2.8.1' currently installed).

```
from pandas.core.computation.check import NUMEXPR_INSTALLED
```

C:\Users\HP\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.4' currently installed).

```
from pandas.core import (
```

C:\Users\HP\anaconda3\lib\site-packages\scipy__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.26.4

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
In [2]: # Load the dataset
df = pd.read_csv("Heart_disease_cleveland_new.csv")
print(df)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	0	145	233	1	2	150	0	2.3	
1	67	1	3	160	286	0	2	108	1	1.5	
2	67	1	3	120	229	0	2	129	1	2.6	
3	37	1	2	130	250	0	0	187	0	3.5	
4	41	0	1	130	204	0	2	172	0	1.4	
..	
298	45	1	0	110	264	0	0	132	0	1.2	
299	68	1	3	144	193	1	0	141	0	3.4	
300	57	1	3	130	131	0	0	115	1	1.2	
301	57	0	1	130	236	0	2	174	0	0.0	
302	38	1	2	138	175	0	0	173	0	0.0	

	slope	ca	thal	target
0	2	0	2	0
1	1	3	1	1
2	1	2	3	1
3	2	0	1	0
4	0	0	1	0
..
298	1	0	3	1
299	1	2	3	1
300	1	1	3	1
301	1	1	1	1
302	0	0	1	0

[303 rows x 14 columns]

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [4]: #managing missing values
missing_values=df.isnull().sum()
print(missing_values)
```


```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```


```
In [5]: df.describe()
```

Out[5]:

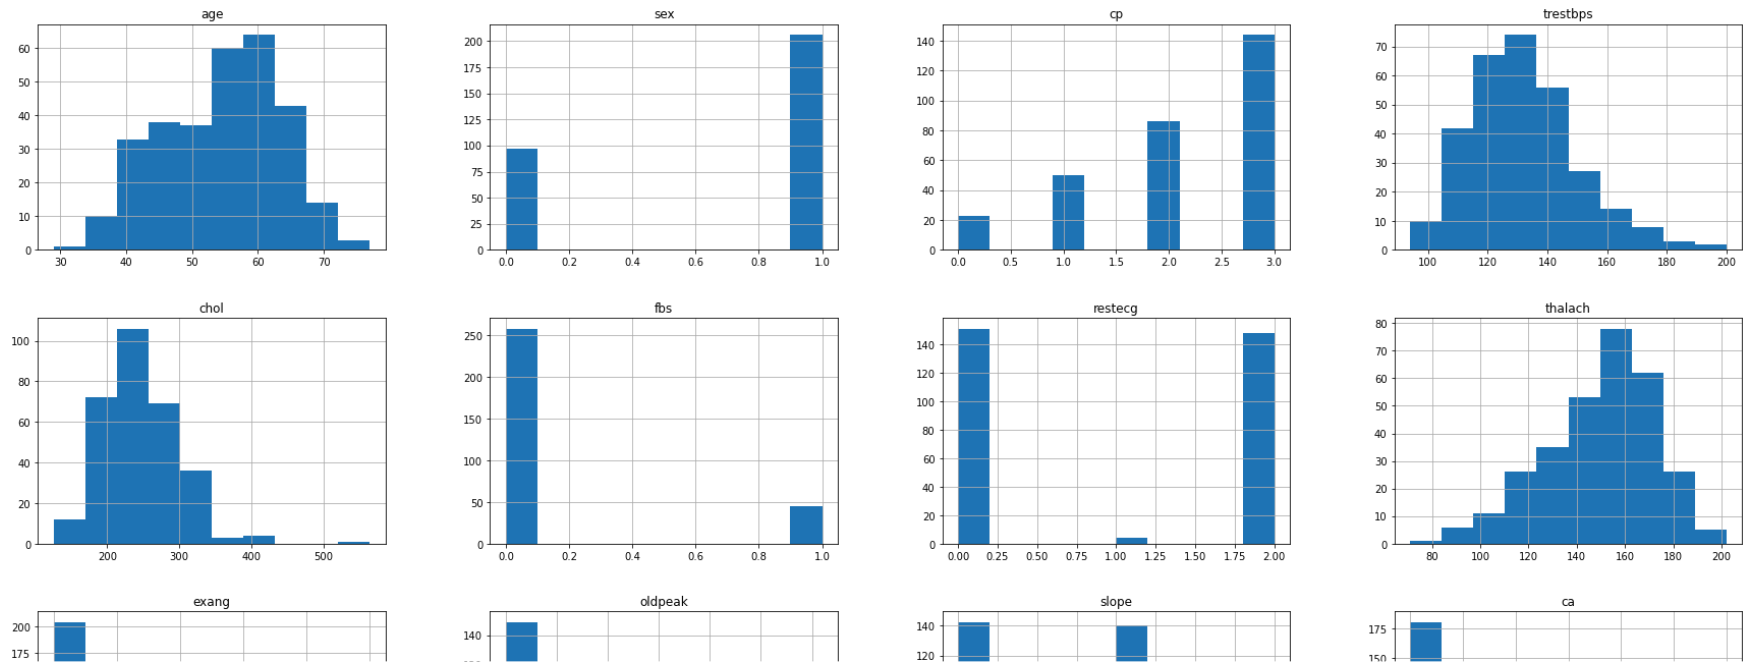
	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.438944	0.679868	2.158416	131.689769	246.693069	0.148515	0.990099	149.607261	0.326733	1.039604	0.600660	0.600660
std	9.038662	0.467299	0.960126	17.599748	51.776918	0.356198	0.994971	22.875003	0.469794	1.161075	0.616226	0.616226
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	48.000000	0.000000	2.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	0.000000	0.000000
50%	56.000000	1.000000	2.000000	130.000000	241.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	1.000000
75%	61.000000	1.000000	3.000000	140.000000	275.000000	0.000000	2.000000	166.000000	1.000000	1.600000	1.000000	1.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	2.000000

```
In [6]: ▶ import seaborn as sns
#get correlations of each features in dataset
corrmat = df.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20),dpi=300)
#plot heat map
g=sns.heatmap(df[top_corr_features].corr(),annot=True,cmap="RdYlGn")
```




```
In [7]: df.hist(figsize=(30, 20))
plt.savefig("histogram.png", dpi=300)
plt.show()
```




```
In [8]: ▶ sns.set_style('whitegrid')  
sns.countplot(x='target',data=df,palette='RdBu_r')
```

Out[8]: <AxesSubplot:xlabel='target', ylabel='count'>



```
In [9]: ▶ # Define the target column  
target_column='target'
```

In [10]:

```
# Separate the dataset based on target values
df_healthy = df[df[target_column] == 0] # No heart disease
df_disease = df[df[target_column] == 1] # Heart disease

# Features to plot (excluding the target column)
features = [col for col in df.columns if col != target_column]


# Create subplots
fig, axes = plt.subplots(1, 2, figsize=(12, 6), dpi=300, sharey=True)

# Convert DataFrame to Long format for Seaborn
df_healthy_melted = df_healthy.melt(value_vars=features, var_name="Feature", value_name="Value")
df_disease_melted = df_disease.melt(value_vars=features, var_name="Feature", value_name="Value")

# Boxplot for target = 0 (No heart disease)
sns.boxplot(y="Feature", x="Value", data=df_healthy_melted, ax=axes[0])
axes[0].set_title("Boxplots of Heart Disease Dataset (Target = 0)")

# Boxplot for target = 1 (Heart disease)
sns.boxplot(y="Feature", x="Value", data=df_disease_melted, ax=axes[1])
axes[1].set_title("Boxplots of Heart Disease Dataset (Target = 1)")

# Adjust layout
plt.tight_layout()
plt.show()
```



```
In [11]: ▶ # Extract features and target variable
X = df.drop(columns=["target"]).values
y = df["target"].values
```

```
In [12]: ▶ # Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
In [13]: ► # Standardize the features  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```



```

In [14]: ▶ # Define the MLPGAN class
class MLPGAN:
    def __init__(self, n_inputs, n_hidden=64, n_outputs=1, population_size=200, generations=100, mutation_rate=0.05,
                 self.n_inputs = n_inputs
                 self.n_hidden = n_hidden
                 self.n_outputs = n_outputs
                 self.dim = (n_inputs * n_hidden) + (n_hidden * n_outputs) + n_hidden + n_outputs
                 self.population_size = population_size
                 self.generations = generations
                 self.mutation_rate = mutation_rate
                 self.crossover_rate = crossover_rate
                 self.population = np.random.randn(self.population_size, self.dim) * 0.01

    def forward_prop(self, params, X):
        input_hidden_weights = params[:self.n_inputs * self.n_hidden].reshape(self.n_inputs, self.n_hidden)
        hidden_output_weights = params[self.n_inputs * self.n_hidden:self.n_inputs * self.n_hidden + self.n_hidden *
        hidden_bias = params[self.n_inputs * self.n_hidden + self.n_hidden * self.n_outputs:self.n_inputs * self.n_hi
        output_bias = params[-self.n_outputs:]

        hidden_layer = np.maximum(0.01 * (np.dot(X, input_hidden_weights) + hidden_bias), np.dot(X, input_hidden_weig
        output_layer = 1 / (1 + np.exp(-(np.dot(hidden_layer, hidden_output_weights) + output_bias)))
        return output_layer

    def fitness_function(self, params, X, y):
        y_pred = self.forward_prop(params, X)
        accuracy = accuracy_score(y, (y_pred >= 0.5).astype(int))
        mse = np.mean((y_pred - y.reshape(-1, 1))**2)
        return accuracy - (0.4 * mse)

    def select_parents(self):
        fitness = np.array([self.fitness_function(ind, X_train_scaled, y_train) for ind in self.population])
        fitness = np.maximum(fitness - fitness.min(), 1e-10)
        probabilities = fitness / fitness.sum()
        selected_indices = np.random.choice(len(probabilities), self.population_size // 2, p=probabilities)
        return self.population[selected_indices]

    def crossover(self, parents):
        offspring = []
        for _ in range(self.population_size - len(parents)):
            if random.random() < self.crossover_rate:
                p1, p2 = random.sample(list(parents), 2)

```

```

        point = random.randint(1, self.dim - 1)
        child = np.concatenate((p1[:point], p2[point:]))
        offspring.append(child)
    return np.array(offspring)

def mutate(self, offspring):
    for i in range(len(offspring)):
        if random.random() < self.mutation_rate:
            mutation_point = random.randint(0, self.dim - 1)
            offspring[i][mutation_point] += np.random.randn() * 0.01
    return offspring

def train(self, X_train, y_train):
    for _ in range(self.generations):
        parents = self.select_parents()
        offspring = self.crossover(parents)
        offspring = self.mutate(offspring)
        self.population = np.vstack((parents, offspring))
        self.best_params = self.select_parents()[-1]

def predict(self, X):
    y_pred = self.forward_prop(self.best_params, X)
    return (y_pred >= 0.5).astype(int)

```

In [15]:  *# Train the genetic algorithm-based MLP model*

```

mlp_ga = MLPGAN(n_inputs=X.shape[1])
mlp_ga.train(X_train_scaled, y_train)

```

In [16]:  *# Generate predictions from MLPGAN*

```

y_pred_mlp_ga = mlp_ga.predict(X_train_scaled)
y_pred_mlp_ga_test = mlp_ga.predict(X_test_scaled)

```

```
In [17]: # Append MLP-GA predictions as additional features
X_train_combined = np.column_stack((X_train_scaled, y_pred_mlp_ga))
X_test_combined = np.column_stack((X_test_scaled, y_pred_mlp_ga_test))
```

```
In [18]: # Train the Random Forest Classifier with tuned parameters
rf = RandomForestClassifier(n_estimators=200, max_depth=10, min_samples_split=4, min_samples_leaf=2, random_state=42)
rf.fit(X_train_combined, y_train)
```

```
Out[18]: RandomForestClassifier(max_depth=10, min_samples_leaf=2, min_samples_split=4,
                                n_estimators=200, random_state=42)
```

```
In [19]: # Make final predictions
y_pred_rf = rf.predict(X_test_combined)
```

```
In [20]: # Compute evaluation metrics
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)
auc_rf = roc_auc_score(y_test, y_pred_rf)
```

```
In [21]: models = {
    "Logistic Regression": LogisticRegression(C=1.5, penalty='l2'),
    "SVM": SVC(C=1, gamma=0.1, kernel='rbf'),
    "KNN": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(criterion='gini'),
    "Random Forest": RandomForestClassifier(n_estimators=1000, criterion='gini'),
    "Extra Trees": ExtraTreesClassifier(n_estimators=100),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100, max_depth=3),
    "GaussianNB": GaussianNB(),
    "XGBoost": XGBClassifier(n_estimators=300, max_depth=15),
    "MLP-BP": MLPClassifier(hidden_layer_sizes=(30,), activation='relu', solver='adam')
}
```



```
In [22]: ▶ results = []
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_test_scaled)[:,-1]
    else:
        y_proba = model.decision_function(X_test_scaled)

    results.append({
        'Model': name,
        'Accuracy': accuracy_score(y_test, y_pred),
        'Precision': precision_score(y_test, y_pred),
        'Recall': recall_score(y_test, y_pred),
        'F1 Score': f1_score(y_test, y_pred),
        'AUC': roc_auc_score(y_test, y_proba)
    })
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
warnings.warn(

```
In [23]: ▶ results.append({
    'Model': 'MLP-GA-RF',
    'Accuracy': accuracy_rf,
    'Precision': precision_rf,
    'Recall': recall_rf,
    'F1 Score': f1_rf,
    'AUC': auc_rf
})
```

```
In [24]: # Print the results
print("MLP+GA+RF Model Metrics:")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision: {precision_rf:.4f}")
print(f"Recall: {recall_rf:.4f}")
print(f"F1 Score: {f1_rf:.4f}")
print(f"AUC: {auc_rf:.4f}")
```

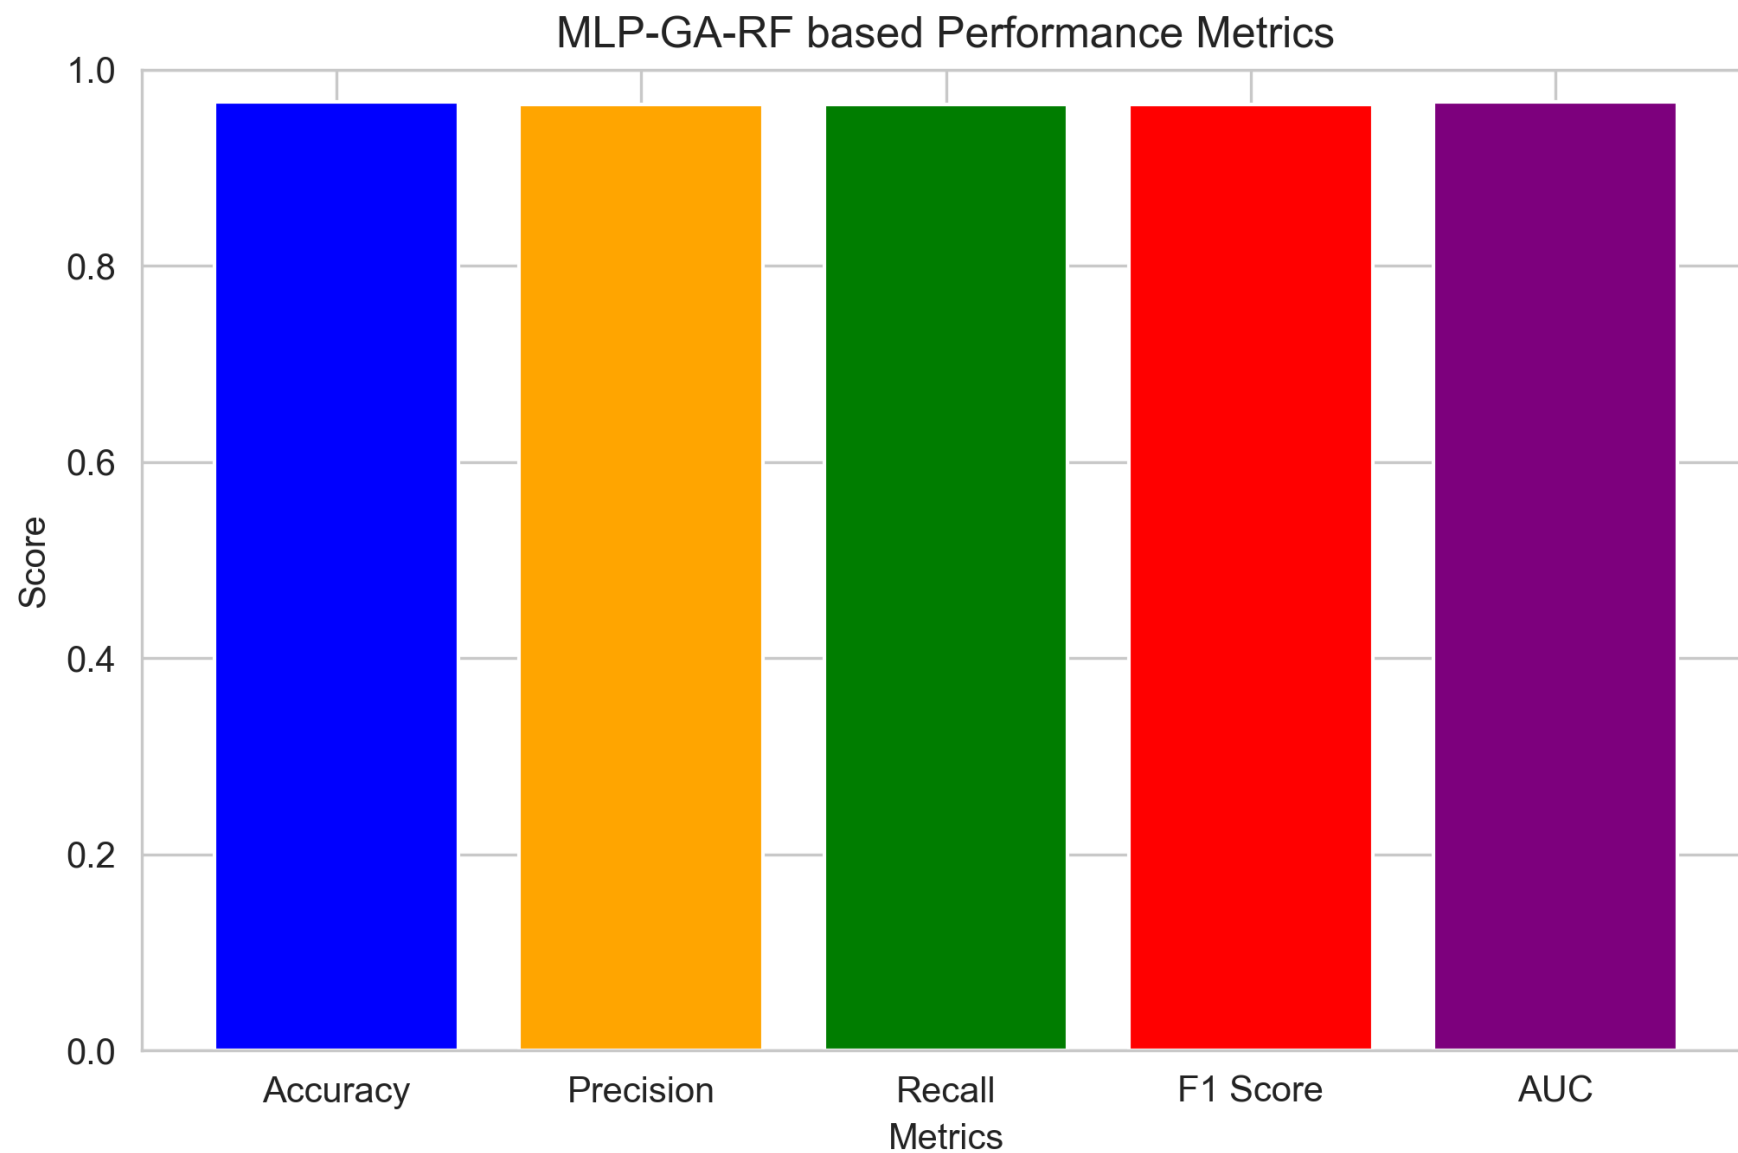
```
MLP+GA+RF Model Metrics:
Accuracy: 0.9672
Precision: 0.9643
Recall: 0.9643
F1 Score: 0.9643
AUC: 0.9670
```

```
In [25]: metrics = {
    "Accuracy": accuracy_rf,
    "Precision": precision_rf,
    "Recall": recall_rf,
    "F1 Score": f1_rf,
    "AUC": auc_rf
}
```

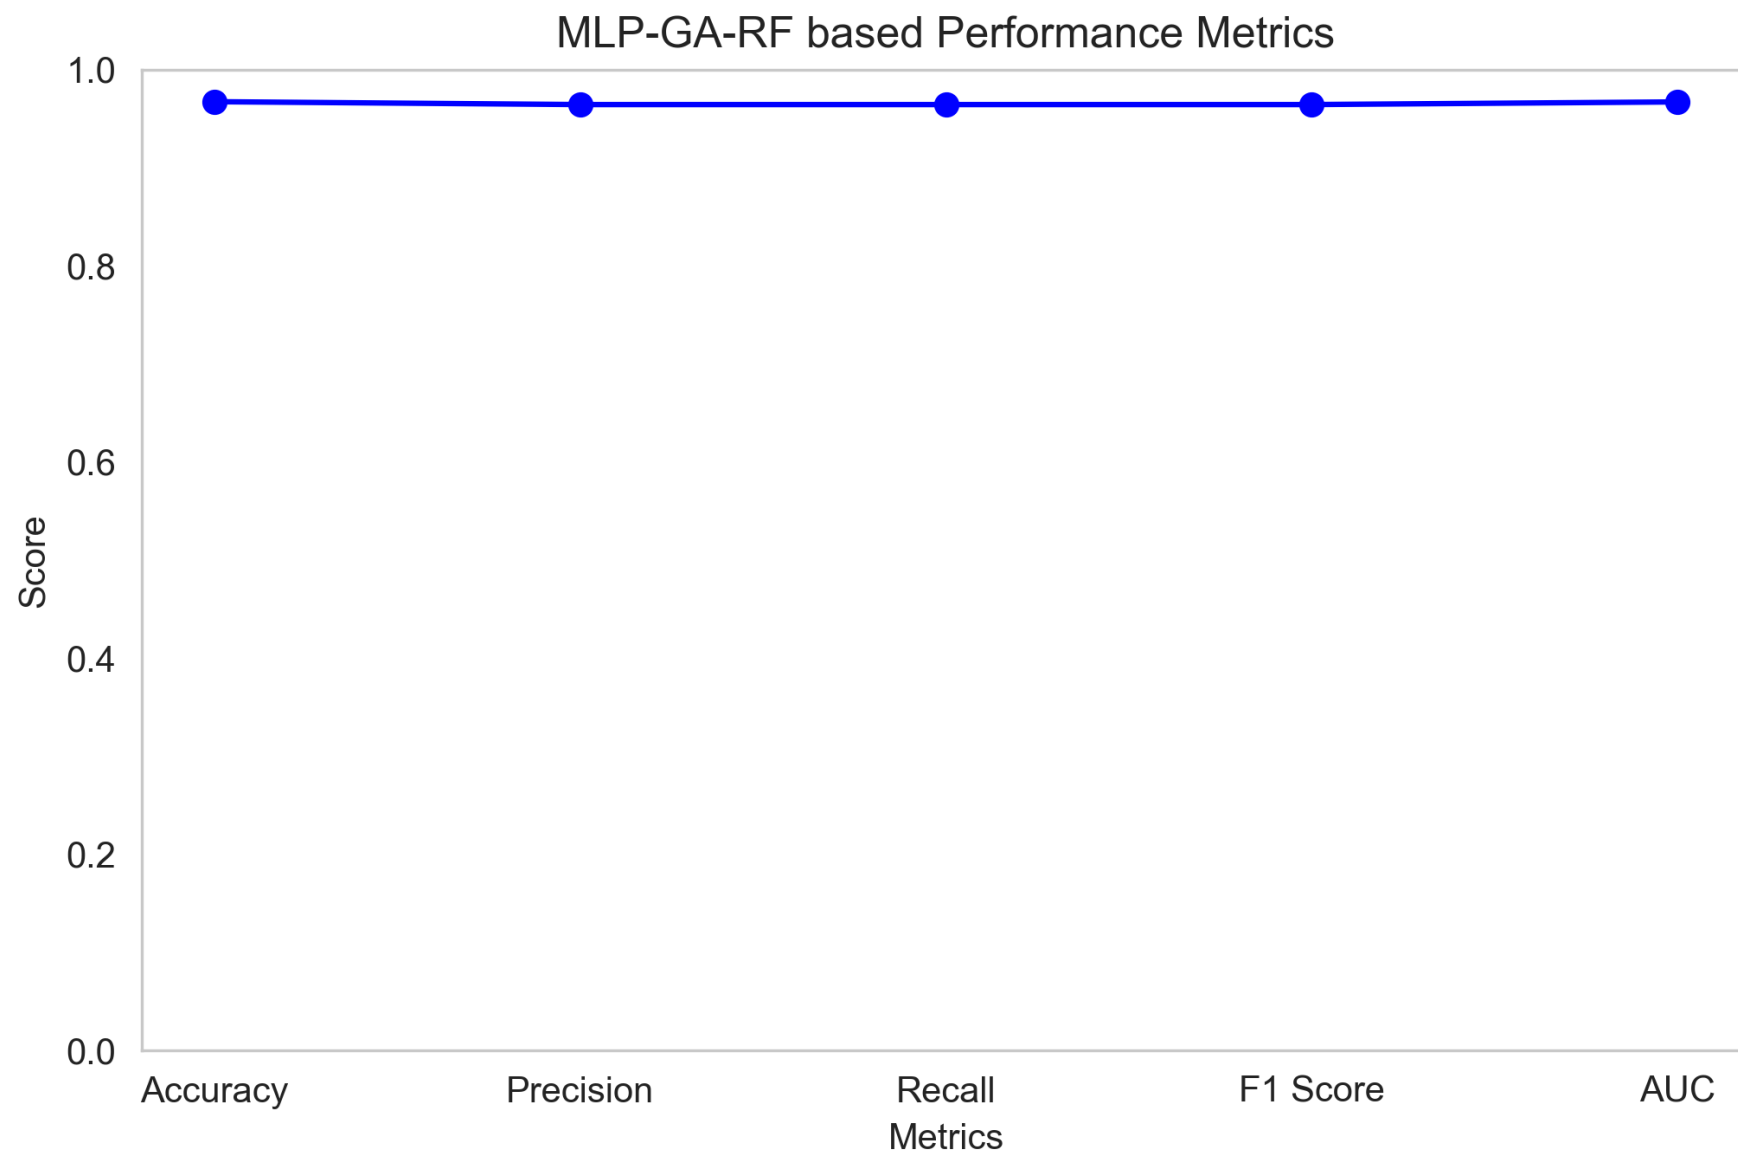
```
In [26]: results_df = pd.DataFrame(results)
print(results_df.sort_values(by='Accuracy', ascending=False))
```

	Model	Accuracy	Precision	Recall	F1 Score	AUC
10	MLP-GA-RF	0.967213	0.964286	0.964286	0.964286	0.966991
2	KNN	0.901639	0.823529	1.000000	0.903226	0.924242
4	Random Forest	0.901639	0.843750	0.964286	0.900000	0.952381
0	Logistic Regression	0.868852	0.812500	0.928571	0.866667	0.952381
7	GaussianNB	0.868852	0.794118	0.964286	0.870968	0.949134
8	XGBoost	0.868852	0.812500	0.928571	0.866667	0.906926
9	MLP-BP	0.868852	0.794118	0.964286	0.870968	0.957792
1	SVM	0.852459	0.806452	0.892857	0.847458	0.944805
6	Gradient Boosting	0.852459	0.787879	0.928571	0.852459	0.945887
5	Extra Trees	0.836066	0.764706	0.928571	0.838710	0.935065
3	Decision Tree	0.786885	0.714286	0.892857	0.793651	0.794913

```
In [27]: ▶ plt.figure(figsize=(8, 5),dpi=300)
plt.bar(metrics.keys(), metrics.values(), color=['blue', 'orange', 'green', 'red', 'purple'])
plt.xlabel("Metrics")
plt.ylabel("Score")
plt.ylim(0, 1)
plt.title("MLP-GA-RF based Performance Metrics")
plt.show()
```




```
In [28]: ▶ plt.figure(figsize=(8, 5),dpi=300)
plt.plot(list(metrics.keys()), list(metrics.values()), marker='o', linestyle='-', color='b')
plt.xlabel("Metrics")
plt.ylabel("Score")
plt.ylim(0, 1)
plt.title("MLP-GA-RF based Performance Metrics")
plt.grid()
plt.show()
```



```
In [29]: ▶ # Confusion matrix visualization
cm = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(5, 4),dpi=300)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No Disease", "Disease"], yticklabels=["No Disease",
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```


Confusion Matrix




NO DISEASE

Predicted

DISEASE

```
In [30]: ▶ # ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_rf)
plt.figure(figsize=(6, 5),dpi=300)
plt.plot(fpr, tpr, color='blue', label=f'ROC curve (AUC = {auc_rf:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for MLP-GA-RF")
#plt.legend()
plt.show()
```


ROC Curve for MLP-GA-RF




False Positive Rate

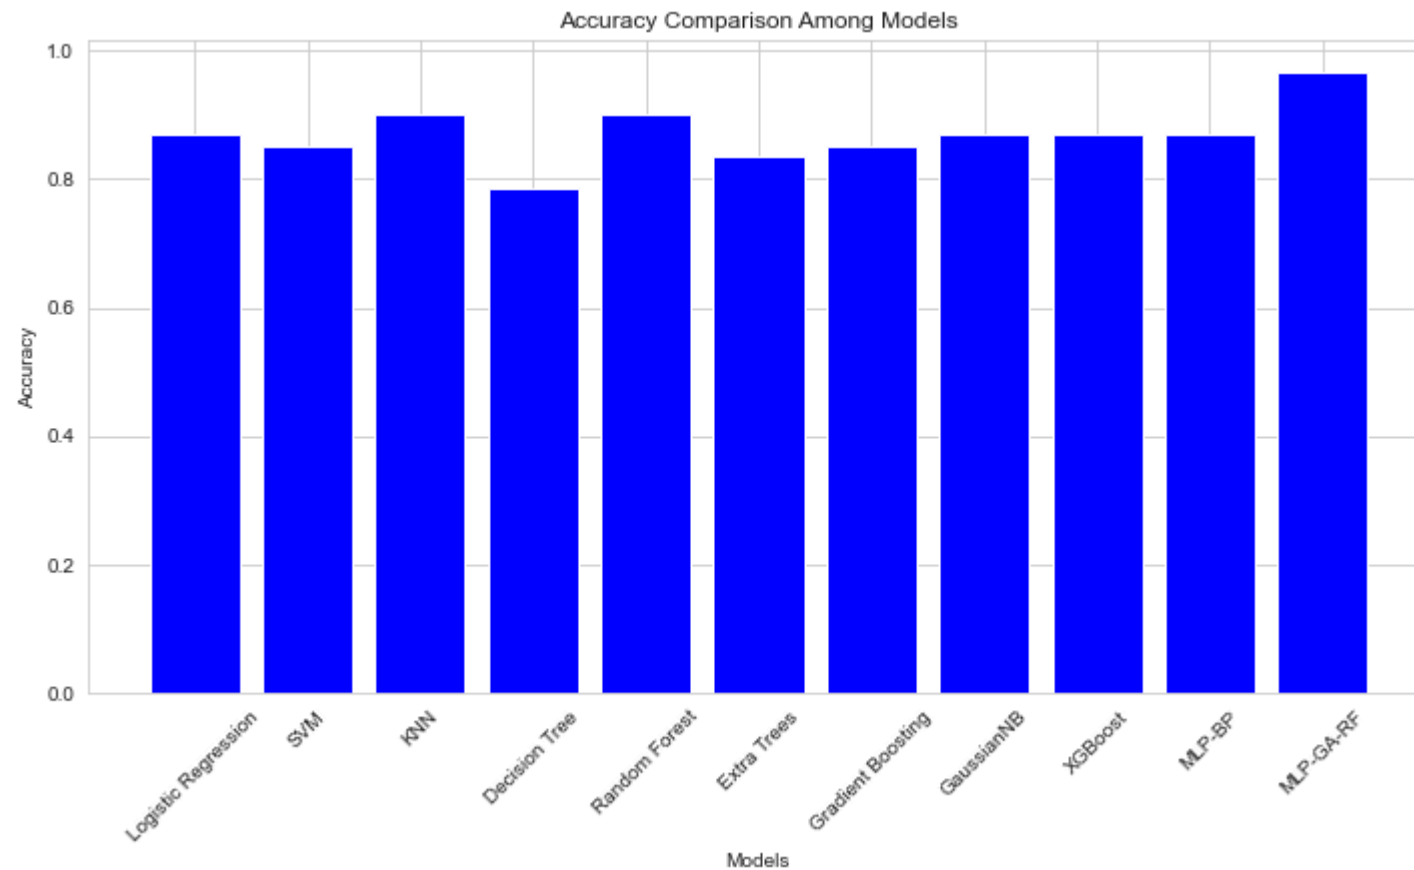
```
In [31]: ▶ # Plot individual ROC curves for each model without AUC and with dpi=300
for name, model in models.items():
    plt.figure(figsize=(5, 5), dpi=300) # High-resolution figure
    model.fit(X_train_scaled, y_train)
    if hasattr(model, "predict_proba"):
        y_probs = model.predict_proba(X_test_scaled)[: , 1] # Get probability scores
    else:
        y_probs = model.decision_function(X_test_scaled) # Use decision_function if predict_proba is unavailable
    fpr, tpr, _ = roc_curve(y_test, y_probs)

    plt.plot(fpr, tpr, linewidth=2, label=f'{name}')
    plt.plot([0, 1], [0, 1], 'k--') # Diagonal Line
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve for {name}')
    plt.legend(loc='lower right')
    plt.show()
```

ROC Curve for Logistic Regression



```
In [32]: ▶ # Bar Plot for Accuracy Comparison
plt.figure(figsize=(12, 6))
plt.bar(results_df['Model'], results_df['Accuracy'], color='blue')
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.title("Accuracy Comparison Among Models")
plt.xticks(rotation=45)
plt.show()
```



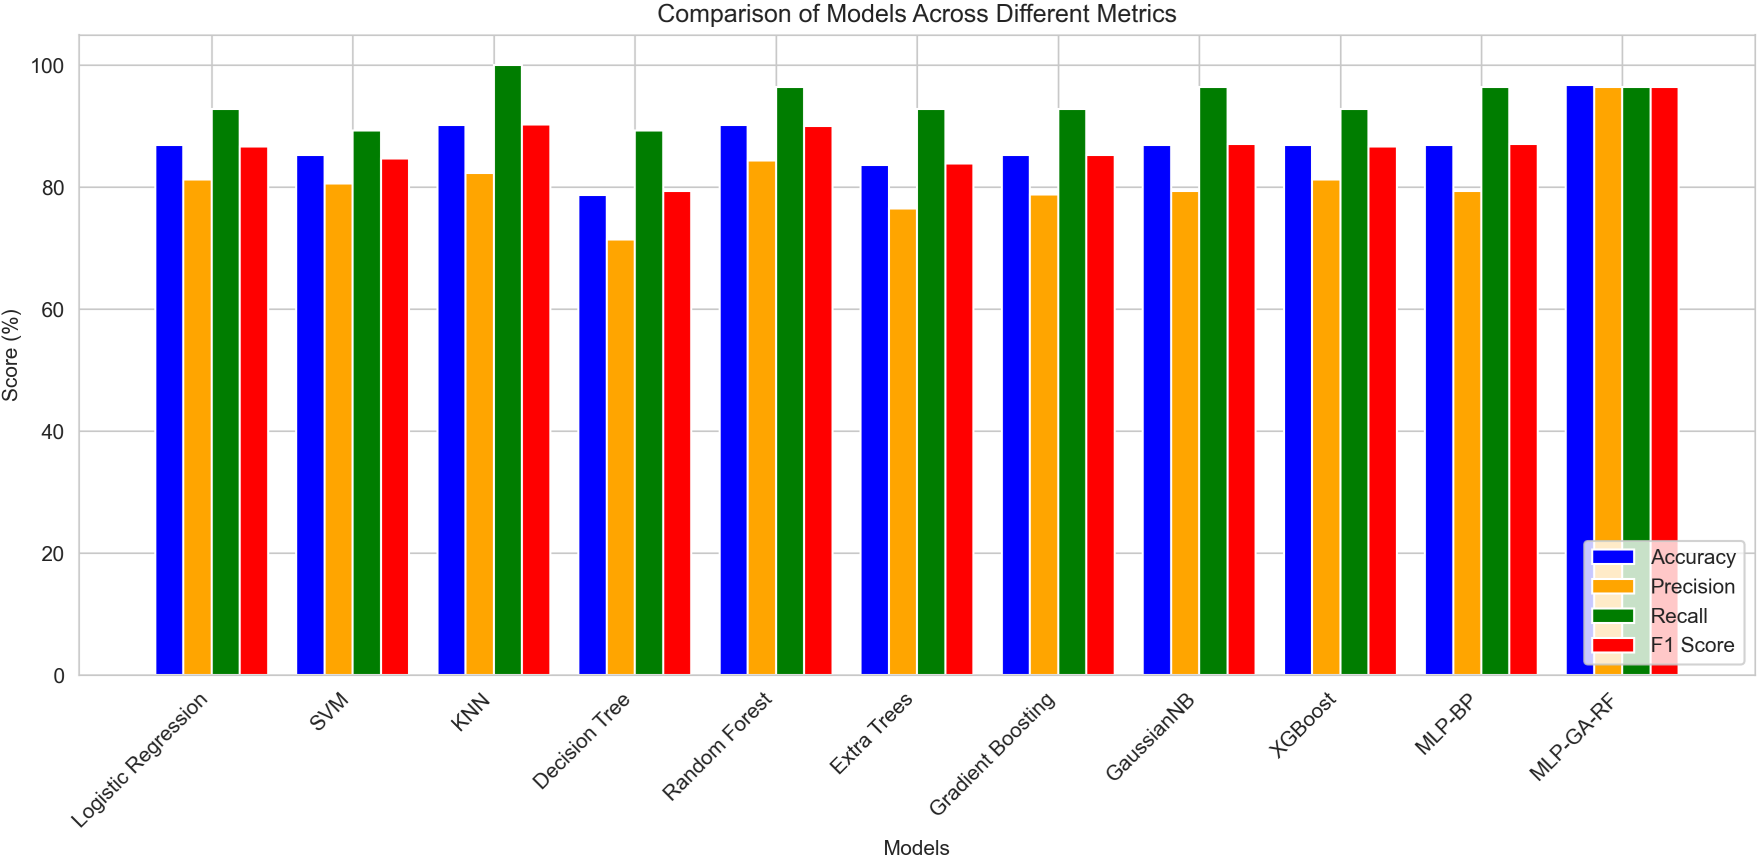
```
In [33]: # Extracting performance metrics for visualization
models_list = results_df['Model'].values
accuracy = results_df['Accuracy'].values * 100
precision = results_df['Precision'].values * 100
recall = results_df['Recall'].values * 100
f1_score_values = results_df['F1 Score'].values * 100

# Set width and positions for bars
x = np.arange(len(models_list))
width = 0.2

# Create a grouped bar chart
fig, ax = plt.subplots(figsize=(12, 6), dpi=300)
ax.bar(x - 1.5 * width, accuracy, width, label="Accuracy", color='blue')
ax.bar(x - 0.5 * width, precision, width, label="Precision", color='orange')
ax.bar(x + 0.5 * width, recall, width, label="Recall", color='green')
ax.bar(x + 1.5 * width, f1_score_values, width, label="F1 Score", color='red')

# Labels and formatting
ax.set_ylabel("Score (%)")
ax.set_xlabel("Models")
ax.set_xticks(x)
ax.set_xticklabels(models_list, rotation=45, ha='right')
ax.set_title("Comparison of Models Across Different Metrics")
ax.legend(loc="lower right")

# Display the plot
plt.tight_layout()
plt.show()
```

In []:

In []:

In []:

```
In [1]: ▶ import numpy as np
import pandas as pd
import random
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix,
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from sklearn.neural_network import MLPClassifier
import pyswarms as ps
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier, GradientBoostingClassifier
from sklearn.metrics import roc_curve, auc
```

C:\Users\HP\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.8.4' or newer of 'numexpr' (version '2.8.1' currently installed).

```
from pandas.core.computation.check import NUMEXPR_INSTALLED
```

C:\Users\HP\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.4' currently installed).

```
from pandas.core import (
```

C:\Users\HP\anaconda3\lib\site-packages\scipy__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.26.4

```
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
In [2]: # Load the dataset
df = pd.read_csv("Heart_disease_cleveland_new.csv")
print(df)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	0	145	233	1	2	150	0	2.3	
1	67	1	3	160	286	0	2	108	1	1.5	
2	67	1	3	120	229	0	2	129	1	2.6	
3	37	1	2	130	250	0	0	187	0	3.5	
4	41	0	1	130	204	0	2	172	0	1.4	
..	
298	45	1	0	110	264	0	0	132	0	1.2	
299	68	1	3	144	193	1	0	141	0	3.4	
300	57	1	3	130	131	0	0	115	1	1.2	
301	57	0	1	130	236	0	2	174	0	0.0	
302	38	1	2	138	175	0	0	173	0	0.0	

	slope	ca	thal	target
0	2	0	2	0
1	1	3	1	1
2	1	2	3	1
3	2	0	1	0
4	0	0	1	0
..
298	1	0	3	1
299	1	2	3	1
300	1	1	3	1
301	1	1	1	1
302	0	0	1	0

[303 rows x 14 columns]

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [4]: #managing missing values
missing_values=df.isnull().sum()
print(missing_values)
```


```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```


```
In [5]: df.describe()
```

Out[5]:

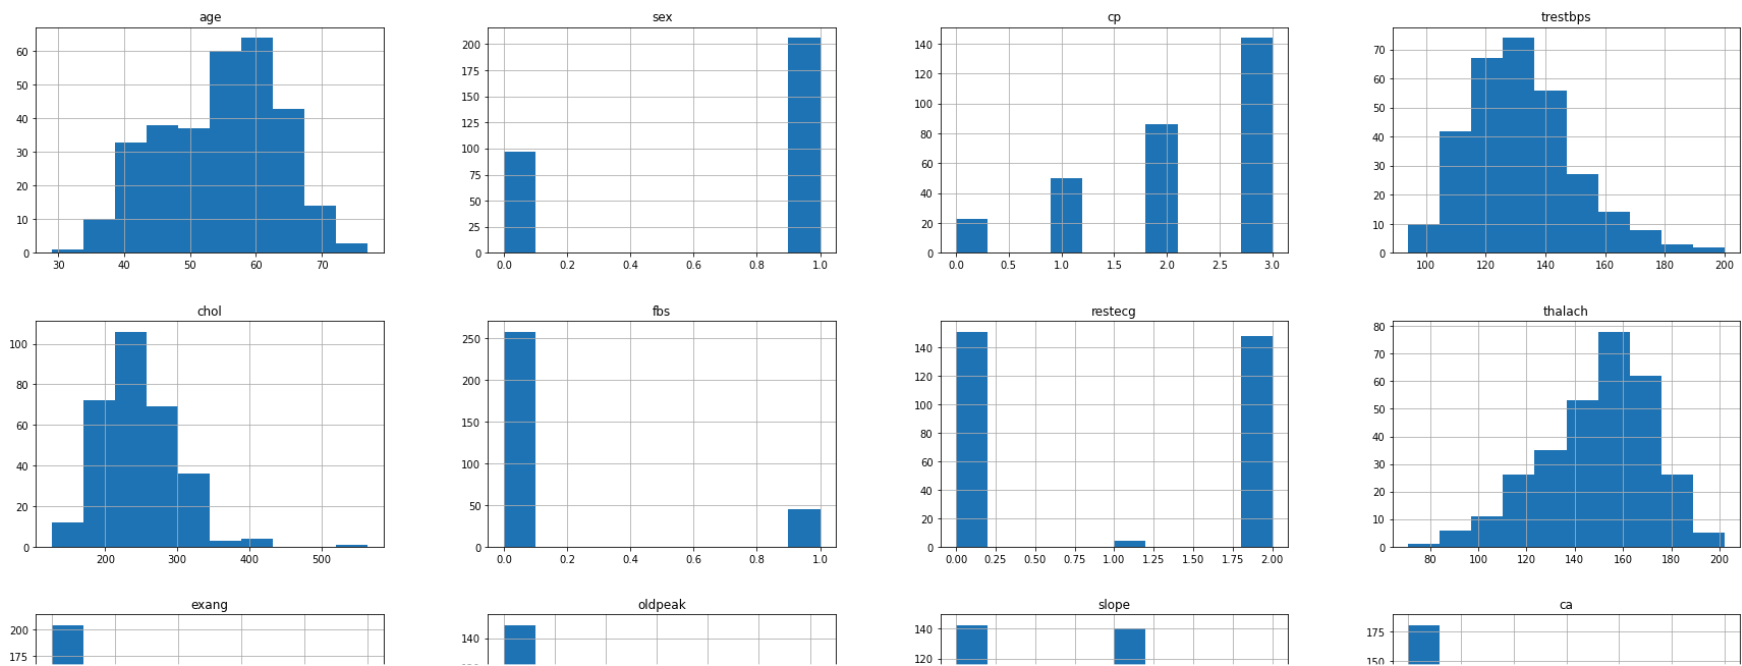
	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303
mean	54.438944	0.679868	2.158416	131.689769	246.693069	0.148515	0.990099	149.607261	0.326733	1.039604	0.600660	C
std	9.038662	0.467299	0.960126	17.599748	51.776918	0.356198	0.994971	22.875003	0.469794	1.161075	0.616226	C
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000	0.000000	C
25%	48.000000	0.000000	2.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000	0.000000	C
50%	56.000000	1.000000	2.000000	130.000000	241.000000	0.000000	1.000000	153.000000	0.000000	0.800000	1.000000	C
75%	61.000000	1.000000	3.000000	140.000000	275.000000	0.000000	2.000000	166.000000	1.000000	1.600000	1.000000	1
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000	2.000000	3

```
In [6]: ▶ import seaborn as sns
#get correlations of each features in dataset
corrmat = df.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20),dpi=300)
#plot heat map
g=sns.heatmap(df[top_corr_features].corr(),annot=True,cmap="RdYlGn")
```




```
In [7]: df.hist(figsize=(30, 20))
plt.savefig("histogram.png", dpi=300)
plt.show()
```



```
In [8]: ▶ sns.set_style('whitegrid')  
sns.countplot(x='target',data=df,palette='RdBu_r')
```

Out[8]: <AxesSubplot:xlabel='target', ylabel='count'>



```
In [9]: ▶ # Define the target column  
target_column='target'
```

In [10]:

```
# Separate the dataset based on target values
df_healthy = df[df[target_column] == 0] # No heart disease
df_disease = df[df[target_column] == 1] # Heart disease

# Features to plot (excluding the target column)
features = [col for col in df.columns if col != target_column]

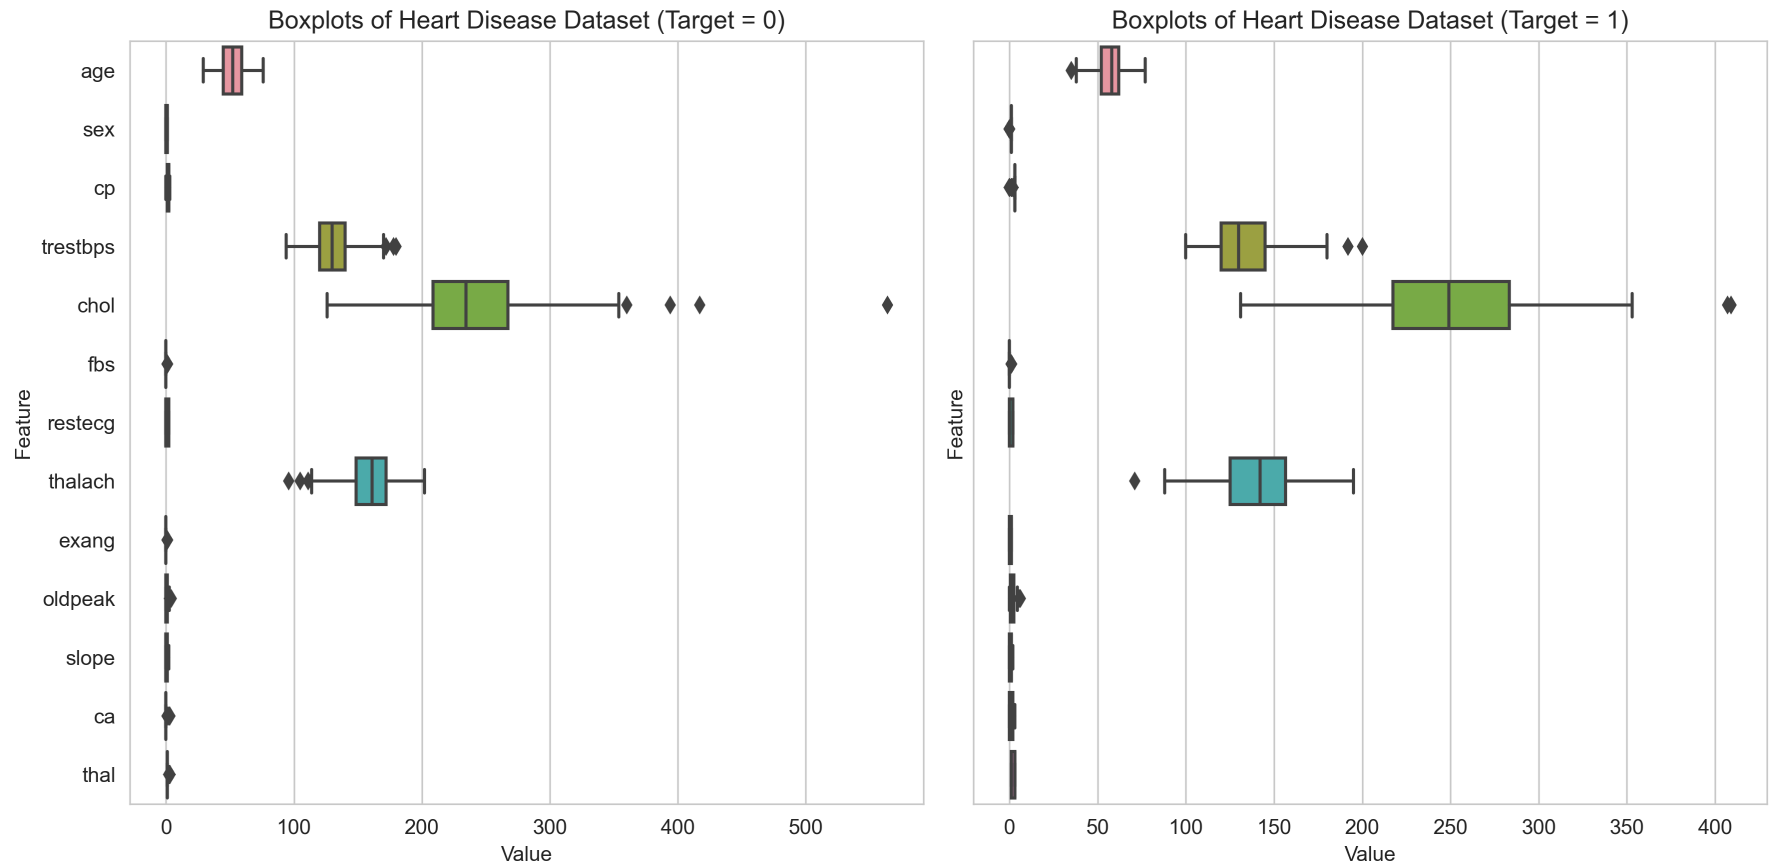
# Create subplots
fig, axes = plt.subplots(1, 2, figsize=(12, 6), dpi=300, sharey=True)

# Convert DataFrame to Long format for Seaborn
df_healthy_melted = df_healthy.melt(value_vars=features, var_name="Feature", value_name="Value")
df_disease_melted = df_disease.melt(value_vars=features, var_name="Feature", value_name="Value")

# Boxplot for target = 0 (No heart disease)
sns.boxplot(y="Feature", x="Value", data=df_healthy_melted, ax=axes[0])
axes[0].set_title("Boxplots of Heart Disease Dataset (Target = 0)")

# Boxplot for target = 1 (Heart disease)
sns.boxplot(y="Feature", x="Value", data=df_disease_melted, ax=axes[1])
axes[1].set_title("Boxplots of Heart Disease Dataset (Target = 1)")

# Adjust layout
plt.tight_layout()
plt.show()
```



```
In [11]: ▶ # Extract features and target variable
X = df.drop(columns=["target"]).values
y = df["target"].values
```

```
In [12]: ▶ # Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
In [13]: ► # Standardize the features  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```



```

In [14]: ▶ # Define the MLPGAN class
class MLPGAN:
    def __init__(self, n_inputs, n_hidden=64, n_outputs=1, population_size=200, generations=100, mutation_rate=0.05,
                 self.n_inputs = n_inputs
                 self.n_hidden = n_hidden
                 self.n_outputs = n_outputs
                 self.dim = (n_inputs * n_hidden) + (n_hidden * n_outputs) + n_hidden + n_outputs
                 self.population_size = population_size
                 self.generations = generations
                 self.mutation_rate = mutation_rate
                 self.crossover_rate = crossover_rate
                 self.population = np.random.randn(self.population_size, self.dim) * 0.01

    def forward_prop(self, params, X):
        input_hidden_weights = params[:self.n_inputs * self.n_hidden].reshape(self.n_inputs, self.n_hidden)
        hidden_output_weights = params[self.n_inputs * self.n_hidden:self.n_inputs * self.n_hidden + self.n_hidden *
        hidden_bias = params[self.n_inputs * self.n_hidden + self.n_hidden * self.n_outputs:self.n_inputs * self.n_hi
        output_bias = params[-self.n_outputs:]

        hidden_layer = np.maximum(0.01 * (np.dot(X, input_hidden_weights) + hidden_bias), np.dot(X, input_hidden_weig
        output_layer = 1 / (1 + np.exp(-(np.dot(hidden_layer, hidden_output_weights) + output_bias)))
        return output_layer

    def fitness_function(self, params, X, y):
        y_pred = self.forward_prop(params, X)
        accuracy = accuracy_score(y, (y_pred >= 0.5).astype(int))
        mse = np.mean((y_pred - y.reshape(-1, 1))**2)
        return accuracy - (0.4 * mse)

    def select_parents(self):
        fitness = np.array([self.fitness_function(ind, X_train_scaled, y_train) for ind in self.population])
        fitness = np.maximum(fitness - fitness.min(), 1e-10)
        probabilities = fitness / fitness.sum()
        selected_indices = np.random.choice(len(probabilities), self.population_size // 2, p=probabilities)
        return self.population[selected_indices]

    def crossover(self, parents):
        offspring = []
        for _ in range(self.population_size - len(parents)):
            if random.random() < self.crossover_rate:
                p1, p2 = random.sample(list(parents), 2)

```

```

        point = random.randint(1, self.dim - 1)
        child = np.concatenate((p1[:point], p2[point:]))
        offspring.append(child)
    return np.array(offspring)

def mutate(self, offspring):
    for i in range(len(offspring)):
        if random.random() < self.mutation_rate:
            mutation_point = random.randint(0, self.dim - 1)
            offspring[i][mutation_point] += np.random.randn() * 0.01
    return offspring

def train(self, X_train, y_train):
    for _ in range(self.generations):
        parents = self.select_parents()
        offspring = self.crossover(parents)
        offspring = self.mutate(offspring)
        self.population = np.vstack((parents, offspring))
        self.best_params = self.select_parents()[-1]

def predict(self, X):
    y_pred = self.forward_prop(self.best_params, X)
    return (y_pred >= 0.5).astype(int)

```

In [15]:  *# Train the genetic algorithm-based MLP model*

```

mlp_ga = MLPGAN(n_inputs=X.shape[1])
mlp_ga.train(X_train_scaled, y_train)

```

In [16]:  *# Generate predictions from MLPGAN*

```

y_pred_mlp_ga = mlp_ga.predict(X_train_scaled)
y_pred_mlp_ga_test = mlp_ga.predict(X_test_scaled)

```



```
In [17]: ▶ # Append MLP-GA predictions as additional features
X_train_combined = np.column_stack((X_train_scaled, y_pred_mlp_ga))
X_test_combined = np.column_stack((X_test_scaled, y_pred_mlp_ga_test))
```

```
In [18]: ▶ # Train the Random Forest Classifier with tuned parameters
rf = RandomForestClassifier(n_estimators=200, max_depth=10, min_samples_split=4, min_samples_leaf=2, random_state=42)
rf.fit(X_train_combined, y_train)
```

```
Out[18]: RandomForestClassifier(max_depth=10, min_samples_leaf=2, min_samples_split=4,
                                n_estimators=200, random_state=42)
```

```
In [19]: ▶ # Make final predictions
y_pred_rf = rf.predict(X_test_combined)
```

```
In [20]: ▶ # Compute evaluation metrics
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)
auc_rf = roc_auc_score(y_test, y_pred_rf)
```

```
In [21]: ▶ models = {
    "Logistic Regression": LogisticRegression(C=1.5, penalty='l2'),
    "SVM": SVC(C=1, gamma=0.1, kernel='rbf'),
    "KNN": KNeighborsClassifier(n_neighbors=5),
    "Decision Tree": DecisionTreeClassifier(criterion='gini'),
    "Random Forest": RandomForestClassifier(n_estimators=1000, criterion='gini'),
    "Extra Trees": ExtraTreesClassifier(n_estimators=100),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100, max_depth=3),
    "GaussianNB": GaussianNB(),
    "XGBoost": XGBClassifier(n_estimators=300, max_depth=15),
    "MLP-BP": MLPClassifier(hidden_layer_sizes=(30,), activation='relu', solver='adam')
}
```

```
In [22]: ▶ results = []
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    if hasattr(model, "predict_proba"):
        y_proba = model.predict_proba(X_test_scaled)[:,-1]
    else:
        y_proba = model.decision_function(X_test_scaled)

    results.append({
        'Model': name,
        'Accuracy': accuracy_score(y_test, y_pred),
        'Precision': precision_score(y_test, y_pred),
        'Recall': recall_score(y_test, y_pred),
        'F1 Score': f1_score(y_test, y_pred),
        'AUC': roc_auc_score(y_test, y_proba)
    })
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\normalization_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

```
warnings.warn(
```

```
In [23]: ▶ results.append({
    'Model': 'MLP-GA-RF',
    'Accuracy': accuracy_rf,
    'Precision': precision_rf,
    'Recall': recall_rf,
    'F1 Score': f1_rf,
    'AUC': auc_rf
})
```

```
In [24]: # Print the results
print("MLP+GA+RF Model Metrics:")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision: {precision_rf:.4f}")
print(f"Recall: {recall_rf:.4f}")
print(f"F1 Score: {f1_rf:.4f}")
print(f"AUC: {auc_rf:.4f}")
```


```
MLP+GA+RF Model Metrics:
Accuracy: 0.9672
Precision: 0.9643
Recall: 0.9643
F1 Score: 0.9643
AUC: 0.9670
```

```
In [25]: metrics = {
    "Accuracy": accuracy_rf,
    "Precision": precision_rf,
    "Recall": recall_rf,
    "F1 Score": f1_rf,
    "AUC": auc_rf
}
```


```
In [26]: results_df = pd.DataFrame(results)
print(results_df.sort_values(by='Accuracy', ascending=False))
```

	Model	Accuracy	Precision	Recall	F1 Score	AUC
10	MLP-GA-RF	0.967213	0.964286	0.964286	0.964286	0.966991
2	KNN	0.901639	0.823529	1.000000	0.903226	0.924242
4	Random Forest	0.901639	0.843750	0.964286	0.900000	0.952381
0	Logistic Regression	0.868852	0.812500	0.928571	0.866667	0.952381
7	GaussianNB	0.868852	0.794118	0.964286	0.870968	0.949134
8	XGBoost	0.868852	0.812500	0.928571	0.866667	0.906926
9	MLP-BP	0.868852	0.794118	0.964286	0.870968	0.957792
1	SVM	0.852459	0.806452	0.892857	0.847458	0.944805
6	Gradient Boosting	0.852459	0.787879	0.928571	0.852459	0.945887
5	Extra Trees	0.836066	0.764706	0.928571	0.838710	0.935065
3	Decision Tree	0.786885	0.714286	0.892857	0.793651	0.794913

```
In [27]: ▶ plt.figure(figsize=(8, 5),dpi=300)
plt.bar(metrics.keys(), metrics.values(), color=['blue', 'orange', 'green', 'red', 'purple'])
plt.xlabel("Metrics")
plt.ylabel("Score")
plt.ylim(0, 1)
plt.title("MLP-GA-RF based Performance Metrics")
plt.show()
```




```
In [28]: ▶ plt.figure(figsize=(8, 5),dpi=300)
plt.plot(list(metrics.keys()), list(metrics.values()), marker='o', linestyle='-', color='b')
plt.xlabel("Metrics")
plt.ylabel("Score")
plt.ylim(0, 1)
plt.title("MLP-GA-RF based Performance Metrics")
plt.grid()
plt.show()
```



```
In [29]: ▶ # Confusion matrix visualization
cm = confusion_matrix(y_test, y_pred_rf)
plt.figure(figsize=(5, 4),dpi=300)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["No Disease", "Disease"], yticklabels=["No Disease",
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```


Confusion Matrix




NO DISEASE

Predicted

DISEASE

```
In [30]: ▶ # ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_rf)
plt.figure(figsize=(6, 5),dpi=300)
plt.plot(fpr, tpr, color='blue', label=f'ROC curve (AUC = {auc_rf:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for MLP-GA-RF")
#plt.legend()
plt.show()
```


ROC Curve for MLP-GA-RF




False Positive Rate

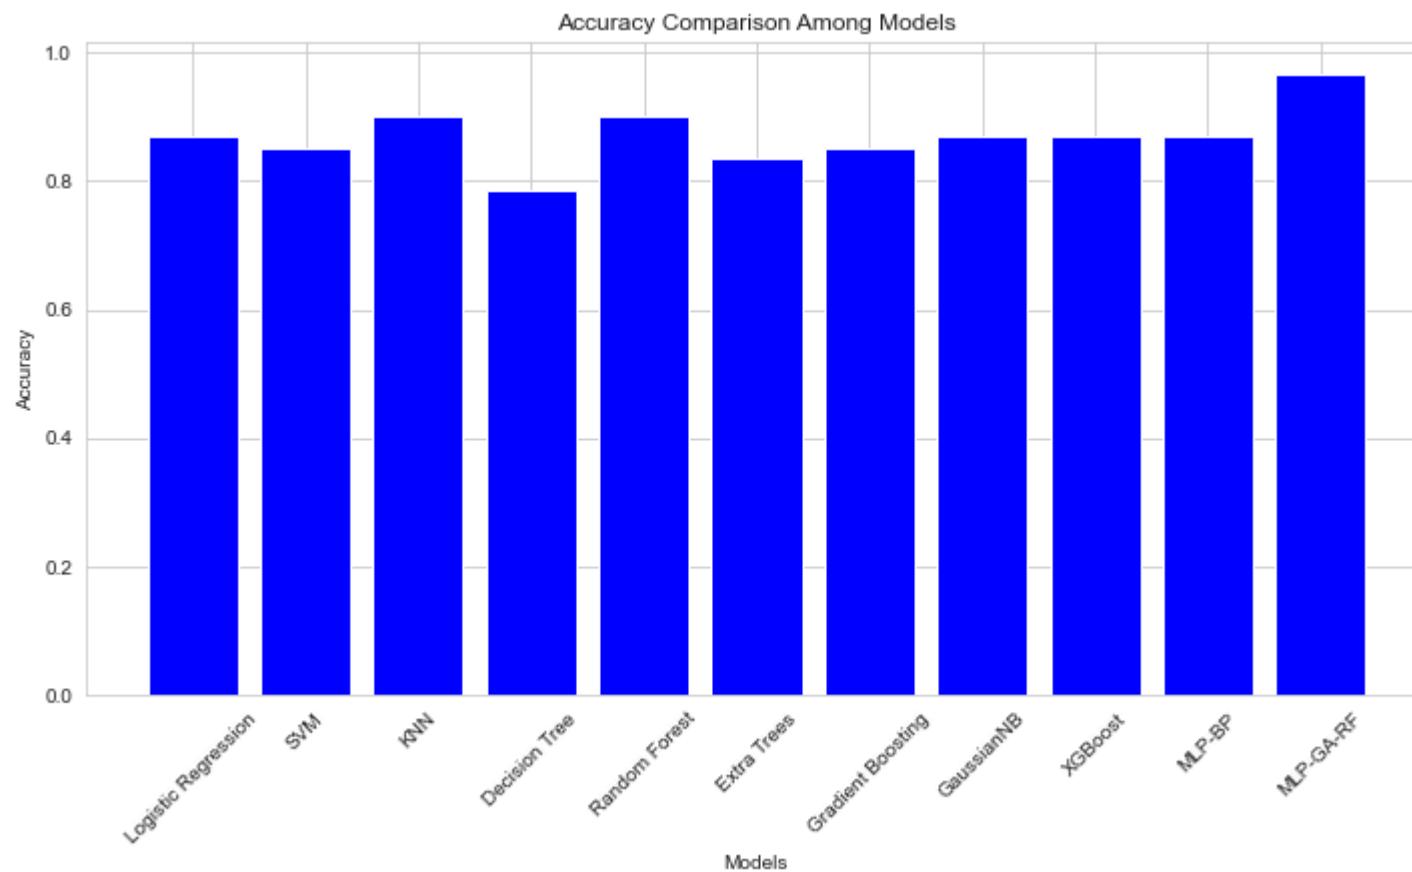
```
In [31]: ▶ # Plot individual ROC curves for each model without AUC and with dpi=300
for name, model in models.items():
    plt.figure(figsize=(5, 5), dpi=300) # High-resolution figure
    model.fit(X_train_scaled, y_train)
    if hasattr(model, "predict_proba"):
        y_probs = model.predict_proba(X_test_scaled)[: , 1] # Get probability scores
    else:
        y_probs = model.decision_function(X_test_scaled) # Use decision_function if predict_proba is unavailable
    fpr, tpr, _ = roc_curve(y_test, y_probs)

    plt.plot(fpr, tpr, linewidth=2, label=f'{name}')
    plt.plot([0, 1], [0, 1], 'k--') # Diagonal Line
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve for {name}')
    plt.legend(loc='lower right')
    plt.show()
```

ROC Curve for Logistic Regression



```
In [32]: ▶ # Bar Plot for Accuracy Comparison
plt.figure(figsize=(12, 6))
plt.bar(results_df['Model'], results_df['Accuracy'], color='blue')
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.title("Accuracy Comparison Among Models")
plt.xticks(rotation=45)
plt.show()
```





```
In [33]: # Extracting performance metrics for visualization
models_list = results_df['Model'].values
accuracy = results_df['Accuracy'].values * 100
precision = results_df['Precision'].values * 100
recall = results_df['Recall'].values * 100
f1_score_values = results_df['F1 Score'].values * 100

# Set width and positions for bars
x = np.arange(len(models_list))
width = 0.2

# Create a grouped bar chart
fig, ax = plt.subplots(figsize=(12, 6), dpi=300)
ax.bar(x - 1.5 * width, accuracy, width, label="Accuracy", color='blue')
ax.bar(x - 0.5 * width, precision, width, label="Precision", color='orange')
ax.bar(x + 0.5 * width, recall, width, label="Recall", color='green')
ax.bar(x + 1.5 * width, f1_score_values, width, label="F1 Score", color='red')

# Labels and formatting
ax.set_ylabel("Score (%)")
ax.set_xlabel("Models")
ax.set_xticks(x)
ax.set_xticklabels(models_list, rotation=45, ha='right')
ax.set_title("Comparison of Models Across Different Metrics")
ax.legend(loc="lower right")

# Display the plot
plt.tight_layout()
plt.show()
```



In []:

In []:

In []:

